

Lab 5: Implementation of Basic CPU Scheduling

CS363 • Operating System (Lab)

Date: 08 - 09 - 2025

CPU Scheduling in Operating System

CPU scheduling is a process that allows one process to use the CPU while the execution of another process is on hold (in waiting state) due to unavailability of any resource like I/O etc, thereby making full use of CPU. The aim of CPU scheduling is to make the system efficient, fast, and fair.

Whenever the CPU becomes idle, the operating system must select one of the processes in the ready queue to be executed. The selection process is carried out by the short-term scheduler (or CPU scheduler). The scheduler selects from among the processes in memory that are ready to execute and allocates the CPU to one of them.

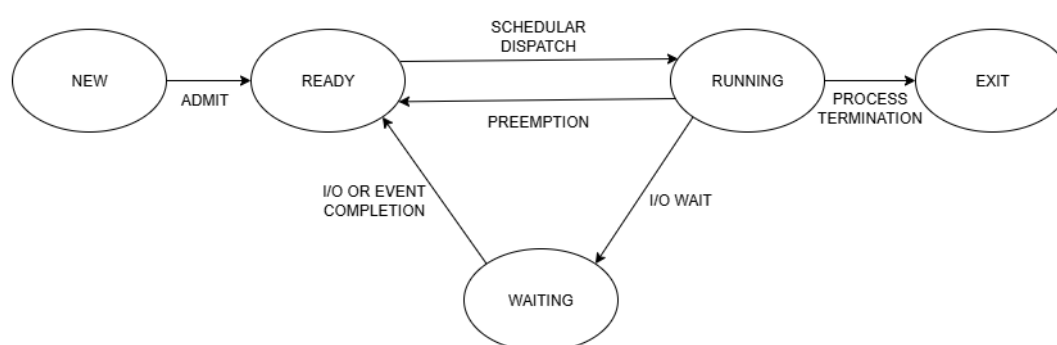


Figure 1: Process state transitions and where scheduling decisions occur.

Types of CPU Scheduling

CPU scheduling decisions may take place under the following four circumstances:

1. When a process switches from the running state to the waiting state (for I/O request or invocation of wait for the termination of one of the child processes).
2. When a process switches from the running state to the ready state (for example, when an interrupt occurs).
3. When a process switches from the waiting state to the ready state (for example, completion of I/O).
4. When a process terminates.

In circumstances 1 and 4, there is no choice in terms of scheduling. A new process (if one exists in the ready queue) must be selected for execution. There is a choice, however, in circumstances 2 and 3.

When scheduling takes place only under circumstances 1 and 4, we say the scheduling scheme is non-preemptive; otherwise, the scheduling scheme is preemptive.

Non-Preemptive Scheduling

Under non-preemptive scheduling, once the CPU is allocated to a process, that process keeps the CPU until it **voluntarily releases** it by either terminating or blocking for an event/I/O (switching to the waiting state). Only then does the scheduler select the next process from the ready queue.

In non-preemptive schemes, the OS does *not* forcibly interrupt a running process just because another process becomes ready or an I/O completion occurs. Device interrupts are still handled, but after the brief interrupt handler finishes, control returns to the *same* running process.

Common non-preemptive algorithms include: **First-Come, First-Served (FCFS)**, **Shortest Job First (SJF, non-preemptive)**, and **Priority (non-preemptive)**.

Preemptive Scheduling

In this type of scheduling, the tasks are usually assigned with priorities. At times it is necessary to run a certain task that has a higher priority before another task although it is running. Therefore, the running task is interrupted for some time and resumed later when the priority task has finished its execution.

Thus this type of scheduling is used mainly when a process switches either from running state to ready state or from waiting state to ready state. The resources (that is CPU cycles) are mainly allocated to the process for a limited amount of time and then are taken away, and after that, the process is again placed back in the ready queue in the case if that process still has a CPU burst time remaining. That process stays in the ready queue until it gets the next chance to execute.

Some algorithms that are based on **preemptive** scheduling are **Round Robin (RR)**, **Shortest Remaining Time First (SRTF)**, and **Priority (preemptive version)** scheduling, etc.

CPU Scheduling: Scheduling Criteria

There are several standard criteria used to judge a “good” scheduling algorithm:

1. **CPU Utilization**
Keep the CPU as *busy as possible*; aim for high utilization.
2. **Throughput**
Number of processes that *complete* per unit time.
3. **Turnaround Time**
Total time to execute a process from arrival to completion (wall-clock): $TAT = CT - AT$.
4. **Waiting Time**
Total time a process spends in the *ready queue* (excludes CPU and I/O time): $WT = TAT - BT$ (no-I/O model).
5. **Response Time**
Time from request/arrival until the *first response is produced* (i.e., until the process first gets the CPU; not final output). $RT = ST - AT$. *Important in time-sharing systems.*

In general, we **maximize** CPU utilization and throughput, and **minimize** turnaround, waiting, and response times (trade-offs may apply).

Notation: AT = arrival time, ST = first start time, CT = completion time, BT = total CPU burst time.

Scheduling Algorithms

To decide which process to execute first and which process to execute last to achieve maximum CPU utilization, computer scientists have defined some algorithms, they are:

- First Come First Serve (FCFS) Scheduling
- Shortest-Job-First (SJF) Scheduling
- Priority Scheduling
- Round Robin (RR) Scheduling
- Multilevel Queue Scheduling
- Multilevel Feedback Queue Scheduling
- Shortest Remaining Time First (SRTF)
- Longest Remaining Time First (LRTF)

First-Come, First-Served (FCFS) Scheduling

The **First-Come, First-Served (FCFS)** scheduling algorithm is the simplest type of CPU scheduling. It is **non-preemptive**, meaning that once the CPU has been allocated to a process, the process keeps it until it either terminates or voluntarily requests I/O. Processes are scheduled in the order in which they arrive in the ready queue.

The implementation of the FCFS policy is easily managed with a **FIFO (First-In, First-Out) queue**. When a process enters the ready queue, its Process Control Block (PCB) is linked to the tail of the queue. When the CPU becomes free, it is allocated to the process at the head of the queue. The running process is then removed from the queue.




Figure 2: FIFO Queue Representation of FCFS Scheduling

While FCFS is simple and ensures fairness based on arrival order, it suffers from several drawbacks. One major problem is the **convoy effect**, where smaller processes are forced to wait if a long CPU-bound process arrives first. This results in poor utilization and high waiting times for shorter jobs. Moreover, the **average waiting time under FCFS is usually quite high**, especially when there is a large variation in process burst times, making it less suitable for interactive and time-sharing systems.

FCFS Scheduling Program

Listing 1: C Program for FCFS Scheduling

```
1  #include <stdio.h>
2  #include <string.h>
3
4  #define MAXN 200
5  #define MAXSEGS (3*MAXN)
6  #define MAXLINE 200
7
8
9  typedef struct {
10     int id;
11     int at;
12     int bt;
13     int ct;
14     int tat;
15     int wt;
16     int input_idx;
17 } Proc;
18
19 typedef struct {
20     char label[16];
21     int start;
22     int end;
23 } Segment;
24
25 static void stable_sort_by_arrival(Proc p[], int n) {
26     int i, j;
27     for (i = 1; i < n; ++i) {
28         Proc key = p[i];
29         j = i - 1;
30         while (j >= 0 && (p[j].at > key.at ||
31             (p[j].at == key.at && p[j].input_idx > key.input_idx))) {
32             p[j + 1] = p[j];
33             --j;
34         }
35         p[j + 1] = key;
36     }
37 }
38
39 static int write_repeat(char *buf, int pos, int max, char ch, int count) {
40     int k;
41     for (k = 0; k < count && pos < max; ++k) buf[pos++] = ch;
42     return pos;
43 }
44
45 static int write_centered(char *buf, int pos, int max, const char *text,
46     int w) {
47     int len = (int)strlen(text);
48     int i;
49     if (w < 1) return pos;
50     if (len > w - 2) {
51         for (i = 0; i < w && pos < max; ++i) buf[pos++] = ' ';
52         return pos;
53     }
54     int left = (w - len) / 2;
55     for (i = 0; i < left && pos < max; ++i) buf[pos++] = ' ';
```

```

55     for (i = 0; i < len && pos < max; ++i) buf[pos++] = text[i];
56     while ((left + len) < w && pos < max) { buf[pos++] = ' '; ++left; }
57     return pos;
58 }
59
60 static void advance_spaces(int *cursor, int n) {
61     int s;
62     for (s = 0; s < n; ++s) putchar(' ');
63     *cursor += n;
64 }
65
66 static void render_gantt(Segment segs[], int segc, int first_time, int
last_time) {
67     if (segc <= 0 || last_time <= first_time) {
68         printf("(no timeline)\n");
69         return;
70     }
71
72     int total_time = last_time - first_time;
73     int scale = 1;
74     if (total_time <= 60) scale = 2;
75     if (total_time <= 30) scale = 3;
76     if (total_time <= 20) scale = 4;
77     if (scale > 8) scale = 8;
78
79     char line1[4*MAXLINE];
80     int p1 = 0;
81     memset(line1, 0, sizeof(line1));
82
83     printf("\n===== GANTT CHART =====\n");
84
85     {
86         int i;
87         for (i = 0; i < segc; ++i) {
88             int dur = segs[i].end - segs[i].start;
89             int width = dur * scale;
90             if (width < 3) width = 3;
91
92             p1 = write_repeat(line1, p1, (int)sizeof(line1)-1, '|', 1);
93
94             if (strcmp(segs[i].label, "cs") == 0) {
95                 p1 = write_repeat(line1, p1, (int)sizeof(line1)-1, '-',
width);
96             } else if (strcmp(segs[i].label, "idle") == 0) {
97                 p1 = write_repeat(line1, p1, (int)sizeof(line1)-1, '.',
width);
98             } else {
99                 if (width >= 5) {
100                     p1 = write_repeat(line1, p1, (int)sizeof(line1)-1, '=',
1);
101                     p1 = write_centered(line1, p1, (int)sizeof(line1)-1,
segs[i].label, width - 2);
102                     p1 = write_repeat(line1, p1, (int)sizeof(line1)-1, '=',
1);
103                 } else {
104                     p1 = write_repeat(line1, p1, (int)sizeof(line1)-1, '=',
width);
105                 }

```

```

106     }
107 }
108 p1 = write_repeat(line1, p1, (int)sizeof(line1)-1, '|', 1);
109 line1[p1] = '\0';
110 printf("%s\n", line1);
111 }
112
113 {
114     int cursor = 0;
115     int accum_cols = 0;
116     int i;
117
118     for (i = 0; i < segc; ++i) {
119         int dur = segs[i].end - segs[i].start;
120         int width = dur * scale; if (width < 3) width = 3;
121
122         int boundary_col = accum_cols + i;
123         if (boundary_col > cursor) advance_spaces(&cursor, boundary_col
            - cursor);
124
125         {
126             char buf[32]; int len;
127             len = snprintf(buf, sizeof(buf), "%d", segs[i].start);
128             fputs(buf, stdout);
129             cursor += len;
130         }
131         accum_cols += width;
132     }
133     {
134         int final_boundary_col = accum_cols + segc;
135         if (final_boundary_col > cursor) advance_spaces(&cursor,
            final_boundary_col - cursor);
136
137         {
138             char buf[32]; int len;
139             len = snprintf(buf, sizeof(buf), "%d", last_time);
140             fputs(buf, stdout);
141             cursor += len;
142         }
143         putchar('\n');
144     }
145
146     printf("=====\n");
147 }
148
149 int main(void) {
150     int n, overhead;
151     Proc procs[MAXN];
152     Segment segs[MAXSEGS];
153     int segc = 0;
154
155     printf("Enter number of processes: ");
156     if (scanf("%d", &n) != 1 || n <= 0 || n > MAXN) {
157         printf("Invalid n.\n");
158         return 0;
159     }
160
161     printf("Enter context-switch overhead (integer, e.g., 1): ");

```

```
162     if (scanf("%d", &overhead) != 1 || overhead < 0) {
163         printf("Invalid overhead.\n");
164         return 0;
165     }
166
167     {
168         int i;
169         for (i = 0; i < n; ++i) {
170             procs[i].id = i + 1;
171             procs[i].input_idx = i;
172             printf("Enter Arrival Time and Burst Time for P%d (AT BT): ", i
173                 + 1);
174             if (scanf("%d %d", &procs[i].at, &procs[i].bt) != 2) {
175                 printf("Bad input.\n");
176                 return 0;
177             }
178             if (procs[i].bt < 0) {
179                 printf("Burst time must be non-negative.\n");
180                 return 0;
181             }
182         }
183
184         stable_sort_by_arrival(procs, n);
185
186         {
187             int i;
188             int time = 0;
189             long long sumBT = 0, total_overhead = 0;
190
191             int minAT = procs[0].at;
192             if (time < minAT) {
193                 strcpy(segs[segc].label, "idle");
194                 segs[segc].start = time;
195                 segs[segc].end = minAT;
196                 ++segc;
197                 time = minAT;
198             }
199
200             for (i = 0; i < n; ++i) {
201                 if (time < procs[i].at) {
202                     strcpy(segs[segc].label, "idle");
203                     segs[segc].start = time;
204                     segs[segc].end = procs[i].at;
205                     ++segc;
206                     time = procs[i].at;
207                 }
208
209                 Segment s;
210                 sprintf(s.label, "P%d", procs[i].id);
211                 s.start = time;
212                 s.end = time + procs[i].bt;
213                 sumBT += procs[i].bt;
214                 procs[i].ct = s.end;
215                 segs[segc++] = s;
216                 time = s.end;
217
218                 if (i < n - 1) {
```

```

219         if (time >= procs[i + 1].at && overhead > 0) {
220             Segment cso;
221             strcpy(cso.label, "cs");
222             cso.start = time;
223             cso.end = time + overhead;
224             segs[segc++] = cso;
225             time += overhead;
226             total_overhead += overhead;
227         }
228     }
229 }
230
231 {
232     int j;
233     long long sumWT = 0, sumTAT = 0;
234     for (j = 0; j < n; ++j) {
235         procs[j].tat = procs[j].ct - procs[j].at;
236         procs[j].wt = procs[j].tat - procs[j].bt;
237         sumWT += procs[j].wt;
238         sumTAT += procs[j].tat;
239     }
240
241     {
242         int first_arrival = procs[0].at;
243         for (j = 1; j < n; ++j)
244             if (procs[j].at < first_arrival) first_arrival = procs[
                j].at;
245         render_gantt(segs, segc, first_arrival, time);
246     }
247
248     printf("\n%-11s %-4s %-4s %-4s %-4s %-11s %-4s %-11s\n",
249           "Process ID", "AT", "BT", "CT", "TAT", "(CT-AT)", "WT",
250           "(TAT-BT)");
251     for (j = 0; j < n; ++j) {
252         char pid[16], tat_expr[32], wt_expr[32];
253         snprintf(pid, sizeof(pid), "%d", procs[j].id);
254         snprintf(tat_expr, sizeof(tat_expr), "%d-%d=%d",
255                procs[j].ct, procs[j].at, procs[j].tat);
256         snprintf(wt_expr, sizeof(wt_expr), "%d-%d=%d",
257                procs[j].tat, procs[j].bt, procs[j].wt);
258
259         printf("%-11s %-4d %-4d %-4d %-4d %-11s %-4d %-11s\n",
260                pid,
261                procs[j].at, procs[j].bt, procs[j].ct,
262                procs[j].tat, tat_expr,
263                procs[j].wt, wt_expr);
264     }
265
266     printf("\nAverage TAT = %.2f\n", (double)sumTAT / n);
267     printf("Average WT = %.2f\n", (double)sumWT / n);
268
269     {
270         int k;
271         long long idle_time = 0;
272         for (k = 0; k < segc; ++k) {
273             if (strcmp(segs[k].label, "idle") == 0) {
                int s = segs[k].start < procs[0].at ? procs[0].at :
                    segs[k].start;
            }
        }
    }

```



```

274         int e = segs[k].end;
275         if (e > s) idle_time += (e - s);
276     }
277 }
278 {
279     double total_elapsed = (double)(time - procs[0].at);
280     double efficiency = total_elapsed > 0 ? (double)sumBT /
        total_elapsed * 100.0 : 100.0;
281
282     printf("\nUseful CPU time (sum BT) = %lld\n", sumBT);
283     printf("Total overhead time      = %lld\n",
        total_overhead);
284     printf("Total idle time (>=ATmin)= %lld\n", idle_time);
285     printf("Total elapsed (from ATmin= %d to end= %d) = %.0
        f\n",
286         procs[0].at, time, total_elapsed);
287     printf("Efficiency (Utilization) = %.2f%%\n",
        efficiency);
288 }
289 }
290 }
291 }
292
293     return 0;
294 }

```

Expected Output:

[202463010@paramshavak ~]\$ nano fcfs.c

[202463010@paramshavak ~]\$ gcc -std=c99 -Wall -Wextra -O2 fcfs.c -o fcfs

[202463010@paramshavak ~]\$./fcfs

Enter number of processes: 5

Enter context-switch overhead (integer, e.g., 1): 1

Enter Arrival Time and Burst Time for P1 (AT BT): 4 5

Enter Arrival Time and Burst Time for P2 (AT BT): 6 4

Enter Arrival Time and Burst Time for P3 (AT BT): 0 3

Enter Arrival Time and Burst Time for P4 (AT BT): 6 2

Enter Arrival Time and Burst Time for P5 (AT BT): 5 4

===== GANTT CHART =====

```

|= P3  =|...|= P1  =|---|= P5  =|---|= P2  =|---|= P4  =|
0      3  4      9  10      14  15      19  20      22

```

=====

Process ID	AT	BT	CT	TAT	(CT-AT)	WT	(TAT-BT)
P3	0	3	3	3	3-0=3	0	3-3=0
P1	4	5	9	5	9-4=5	0	5-5=0
P5	5	4	14	9	14-5=9	5	9-4=5
P2	6	4	19	13	19-6=13	9	13-4=9
P4	6	2	22	16	22-6=16	14	16-2=14

Average TAT = 9.20

Average WT = 5.60

Useful CPU time (sum BT) = 18
 Total overhead time = 3
 Total idle time ($\geq AT_{min}$) = 1
 Total elapsed (from $AT_{min} = 0$ to end = 22) = 22
 Efficiency (Utilization) = 81.82%

Shortest-Job-First (SJF) Scheduling

The **Shortest-Job-First (SJF)** scheduling algorithm selects the process with the shortest next CPU burst time for execution. It is considered an optimal scheduling algorithm because it minimizes the average waiting time for a given set of processes. However, the main difficulty is that the length of the next CPU burst is not known in advance, and in practice, it is usually predicted based on the recent history of a process's CPU bursts.

Non-Preemptive SJF

In **Non-Preemptive SJF**, once the CPU has been allocated to a process, it cannot be taken away until the process either terminates or moves to the waiting state. The scheduler simply picks the process with the smallest burst time from the ready queue at the time of allocation. This ensures fairness among short jobs, but longer jobs may suffer from starvation if short jobs keep arriving continuously.

Preemptive SJF (Shortest Remaining Time First, SRTF)

In **Preemptive SJF**, also called **Shortest Remaining Time First (SRTF)**, the CPU can be preempted if a new process arrives with a CPU burst smaller than the remaining time of the currently running process. This leads to better response times in interactive systems, as shorter jobs are quickly executed. However, the overhead of frequent preemptions and the possibility of starvation for longer processes remain key drawbacks.

SJF Scheduling Program

Listing 2: C Program for SJF Scheduling (Preemptive & Non-Preemptive)

```

1  #include <stdio.h>
2  #include <string.h>
3  #include <limits.h>
4
5
6  #define MAXN 200
7  #define MAXSEGS (4*MAXN)
8  #define MAXLINE 256
9
10 typedef struct {
11     int id;
12     int at, bt;
13     int ct, tat, wt;
14     int rem;
15     int done;
16 } Proc;
17
18 typedef struct {
19     char label[16];
20     int start, end;

```

```

21 } Segment;
22
23 static int write_repeat(char *buf, int pos, int max, char ch, int count){
24     for (int k = 0; k < count && pos < max; ++k) buf[pos++] = ch;
25     return pos;
26 }
27
28 static int write_centered(char *buf, int pos, int max, const char *txt, int
    w){
29     int len = (int)strlen(txt);
30     if (w < 1) return pos;
31     if (len >= w) {
32         for (int i=0; i<w && pos<max; ++i) buf[pos++] = txt[i];
33         return pos;
34     }
35     int left = (w - len)/2;
36     for (int i=0; i<left && pos<max; ++i) buf[pos++] = ' ';
37     for (int i=0; i<len && pos<max; ++i) buf[pos++] = txt[i];
38     while ((left+len) < w && pos<max){ buf[pos++] = ' '; ++left; }
39     return pos;
40 }
41
42 static void advance_spaces(int *cursor, int n){
43     for (int i=0; i<n; ++i) putchar(' ');
44     *cursor += n;
45 }
46
47 static void render_gantt(Segment segs[], int segc, int first_time, int
    last_time){
48     if (segc <= 0 || last_time <= first_time){
49         printf("(no timeline)\n");
50         return;
51     }
52
53     int total = last_time - first_time;
54     int scale = 1;
55     if (total <= 60) scale = 2;
56     if (total <= 30) scale = 3;
57     if (total <= 20) scale = 4;
58     if (scale > 8) scale = 8;
59
60     char line1[4*MAXLINE]; int p1 = 0;
61     memset(line1, 0, sizeof(line1));
62
63     printf("\n===== GANTT CHART =====\n");
64     for (int i=0; i<segc; ++i){
65         int dur = segs[i].end - segs[i].start;
66         int w = dur*scale; if (w < 3) w = 3;
67         p1 = write_repeat(line1, p1, (int)sizeof(line1)-1, '|', 1);
68
69         if (strcmp(segs[i].label, "idle")==0){
70             p1 = write_repeat(line1, p1, (int)sizeof(line1)-1, '.', w);
71         }else{
72             if (w >= 5){
73                 p1 = write_repeat(line1, p1, (int)sizeof(line1)-1, '=', 1);
74                 p1 = write_centered(line1, p1, (int)sizeof(line1)-1, segs[i]
                    ].label, w-2);
75                 p1 = write_repeat(line1, p1, (int)sizeof(line1)-1, '=', 1);

```

```

76         }else{
77             p1 = write_centered(line1, p1, (int)sizeof(line1)-1, segs[i]
                ].label, w);
78         }
79     }
80 }
81 p1 = write_repeat(line1, p1, (int)sizeof(line1)-1, '|', 1);
82 line1[p1] = '\0';
83 printf("%s\n", line1);
84
85 int cursor = 0, accum = 0;
86 for (int i=0; i<segc; ++i){
87     int dur = segs[i].end - segs[i].start;
88     int w = dur*scale; if (w < 3) w = 3;
89     int boundary_col = accum + i;
90     if (boundary_col > cursor) advance_spaces(&cursor, boundary_col -
        cursor);
91     char buf[32]; int len = snprintf(buf, sizeof(buf), "%d", segs[i].
        start);
92     fputs(buf, stdout); cursor += len;
93     accum += w;
94 }
95 int final_col = accum + segc;
96 if (final_col > cursor) advance_spaces(&cursor, final_col - cursor);
97 { char buf[32]; int len = snprintf(buf, sizeof(buf), "%d", last_time);
98     fputs(buf, stdout); cursor += len; }
99 putchar('\n');
100 printf("=====\n");
101 }
102
103 static void push_segment(Segment segs[], int *segc, const char *label, int
    s, int e){
104     if (s >= e) return;
105     if (*segc > 0){
106         Segment *prev = &segs[*segc - 1];
107         if (strcmp(prev->label, label)==0 && prev->end == s){
108             prev->end = e; return;
109         }
110     }
111     strncpy(segs[*segc].label, label, sizeof(segs[*segc].label)-1);
112     segs[*segc].label[sizeof(segs[*segc].label)-1] = '\0';
113     segs[*segc].start = s; segs[*segc].end = e;
114     (*segc)++;
115 }
116
117 static void print_table(Proc p[], int n){
118     double avgTAT=0, avgWT=0;
119     printf("\n%-10s %-4s %-4s %-4s %-4s %-10s %-4s %-10s\n",
120         "ProcessID", "AT", "BT", "CT", "TAT", "(CT-AT)", "WT", "(TAT-BT)");
121     for (int i=0; i<n; ++i){
122         char pid[16], tatExpr[32], wtExpr[32];
123         snprintf(pid, sizeof(pid), "%d", p[i].id);
124         snprintf(tatExpr, sizeof(tatExpr), "%d-%d=%d", p[i].ct, p[i].at, p[i].
            tat);
125         snprintf(wtExpr, sizeof(wtExpr), "%d-%d=%d", p[i].tat, p[i].bt, p[i].wt)
            ;
126         avgTAT += p[i].tat; avgWT += p[i].wt;
127         printf("%-10s %-4d %-4d %-4d %-4d %-10s %-4d %-10s\n",

```

```

128         pid, p[i].at, p[i].bt, p[i].ct, p[i].tat,
129         tatExpr, p[i].wt, wtExpr);
130     }
131     printf("\nAverage TAT = %.2f\n", avgTAT/n);
132     printf("Average WT = %.2f\n", avgWT/n);
133 }
134
135 static void sjf_nonpreemptive(Proc p[], int n){
136     int done=0, t=0, first_arr=INT_MAX, last_ct=0;
137     long long sumBT=0;
138     Segment segs[MAXSEGS]; int segc=0;
139     for (int i=0;i<n;++i){ if(p[i].at<first_arr)first_arr=p[i].at; sumBT+=p
        [i].bt; p[i].done=0; }
140     if (t<first_arr) t=first_arr;
141
142     while(done<n){
143         int idx=-1,minBT=INT_MAX;
144         for(int i=0;i<n;++i){
145             if(!p[i].done && p[i].at<=t){
146                 if(p[i].bt<minBT || (p[i].bt==minBT && p[i].at<p[idx].at)){
147                     minBT=p[i].bt; idx=i;
148                 }
149             }
150         }
151         if(idx!=-1){
152             int nextAT=INT_MAX;
153             for(int i=0;i<n;++i) if(!p[i].done && p[i].at>t && p[i].at<
                nextAT) nextAT=p[i].at;
154             push_segment(segs,&segc,"idle",t,nextAT);
155             t=nextAT; continue;
156         }
157         char label[16]; snprintf(label,sizeof(label),"P%d",p[idx].id);
158         push_segment(segs,&segc,label,t,t+p[idx].bt);
159         t+=p[idx].bt; p[idx].ct=t; p[idx].done=1; last_ct=t; done++;
160     }
161
162     for(int i=0;i<n;++i){ p[i].tat=p[i].ct-p[i].at; p[i].wt=p[i].tat-p[i].
        bt; }
163     render_gantt(segs,segc,first_arr,last_ct);
164     print_table(p,n);
165     double total_elapsed=(double)(last_ct-first_arr);
166     double efficiency=total_elapsed>0?(double)sumBT/total_elapsed
        *100.0:100.0;
167     printf("\nUseful CPU time (sum BT) = %lld\n",sumBT);
168     printf("Total elapsed (from ATmin= %d to end= %d) = %.0f\n",first_arr,
        last_ct,total_elapsed);
169     printf("Efficiency (Utilization) = %.2f%%\n",efficiency);
170 }
171
172 static void sjf_preemptive(Proc p[], int n){
173     int done=0,t=0,first_arr=INT_MAX,last_ct=0;
174     long long sumBT=0;
175     Segment segs[MAXSEGS]; int segc=0;
176     for(int i=0;i<n;++i){p[i].rem=p[i].bt;p[i].done=0;if(p[i].at<first_arr)
        first_arr=p[i].at;sumBT+=p[i].bt;}
177     if(t<first_arr) t=first_arr;
178
179     int current=-1,seg_start=t;

```

```

180 while(done<n){
181     int idx=-1,minR=INT_MAX;
182     for(int i=0;i<n;++i){
183         if(p[i].at<=t && p[i].rem>0){
184             if(p[i].rem<minR || (p[i].rem==minR && p[i].at<p[idx].at)){
185                 minR=p[i].rem; idx=i;
186             }
187         }
188     }
189     if(idx==-1){
190         if(current!=-1){
191             char label[16];snprintf(label,sizeof(label),"P%d",p[current
192             ].id);
193             push_segment(segs,&segc,label,seg_start,t); current=-1;
194         }
195         int nextAT=INT_MAX;
196         for(int i=0;i<n;++i) if(p[i].rem>0&&p[i].at>t&&p[i].at<nextAT)
197             nextAT=p[i].at;
198         push_segment(segs,&segc,"idle",t,nextAT); t=nextAT; seg_start=t
199         ; continue;
200     }
201     if(current!=idx){
202         if(current!=-1){
203             char label[16];snprintf(label,sizeof(label),"P%d",p[current
204             ].id);
205             push_segment(segs,&segc,label,seg_start,t);
206         }
207         current=idx; seg_start=t;
208     }
209     p[current].rem--; t++;
210     if(p[current].rem==0){
211         p[current].ct=t; p[current].done=1; done++;
212         char label[16];snprintf(label,sizeof(label),"P%d",p[current].id
213         );
214         push_segment(segs,&segc,label,seg_start,t);
215         current=-1; seg_start=t; if(t>last_ct)last_ct=t;
216     }
217 }
218
219 for(int i=0;i<n;++i){p[i].tat=p[i].ct-p[i].at;p[i].wt=p[i].tat-p[i].bt;
220     if(p[i].ct>last_ct)last_ct=p[i].ct;}
221 render_gantt(segs,segc,first_arr,last_ct);
222 print_table(p,n);
223 double total_elapsed=(double)(last_ct-first_arr);
224 double efficiency=total_elapsed>0?(double)sumBT/total_elapsed
225     *100.0:100.0;
226 printf("\nUseful CPU time (sum BT) = %lld\n",sumBT);
227 printf("Total elapsed (from ATmin= %d to end= %d) = %.0f\n",first_arr,
228     last_ct,total_elapsed);
229 printf("Efficiency (Utilization) = %.2f%%\n",efficiency);
230 }
231
232 int main(void){
233     int n,choice; Proc p[MAXN];
234     printf("Select Scheduling Type:\n1. Preemptive SJF (SRTF)\n2. Non-
235     Preemptive SJF\nChoice: ");
236     if(scanf("%d",&choice)!=1||((choice!=1&&choice!=2))){printf("Invalid
237     choice.\n");return 0;}

```

```

228     printf("Enter number of processes: "); if(scanf("%d",&n)!=1||n<=0||n>
        MAXN){printf("Invalid n.\n");return 0;}
229     for(int i=0;i<n;++i){p[i].id=i+1;printf("Enter Arrival Time and Burst
        Time for P%d (AT BT): ",i+1);
230         if(scanf("%d %d",&p[i].at,&p[i].bt)!=2||p[i].bt<0){printf("Bad
            input.\n");return 0;}
231         p[i].ct=p[i].tat=p[i].wt=0;p[i].rem=p[i].bt;p[i].done=0;}
232     if(choice==2) sjf_nonpreemptive(p,n); else sjf_preemptive(p,n);
233     return 0;
234 }

```

Expected Output (Non-Preemptive):

```

[202463010@paramshavak ~]$ nano sjf4.c
[202463010@paramshavak ~]$ gcc -std=c99 -Wall -Wextra -O2 sjf4.c -o sjf4
[202463010@paramshavak ~]$ ./sjf4

```

Select Scheduling Type:

1. Preemptive SJF (SRTF)
2. Non-Preemptive SJF

Choice: 2

Enter number of processes: 4

Enter Arrival Time and Burst Time for P1 (AT BT): 0 6

Enter Arrival Time and Burst Time for P2 (AT BT): 0 8

Enter Arrival Time and Burst Time for P3 (AT BT): 0 7

Enter Arrival Time and Burst Time for P4 (AT BT): 0 3

```

===== GANTT CHART =====
|= P4  |=      P1      |=      P3      |=      P2      |=
0       3           9           16           24
=====

```

ProcessID	AT	BT	CT	TAT	(CT-AT)	WT	(TAT-BT)
P1	0	6	9	9	9-0=9	3	9-6=3
P2	0	8	24	24	24-0=24	16	24-8=16
P3	0	7	16	16	16-0=16	9	16-7=9
P4	0	3	3	3	3-0=3	0	3-3=0

Average TAT = 13.00

Average WT = 7.00

Useful CPU time (sum BT) = 24

Total elapsed (from ATmin= 0 to end= 24) = 24

Efficiency (Utilization) = 100.00%

Expected Output (Preemptive):

```

[202463010@paramshavak ~]$ nano sjf4.c
[202463010@paramshavak ~]$ gcc -std=c99 -Wall -Wextra -O2 sjf4.c -o sjf4
[202463010@paramshavak ~]$ ./sjf4

```

Select Scheduling Type:

1. Preemptive SJF (SRTF)
2. Non-Preemptive SJF

Choice: 1

Enter number of processes: 4

Enter Arrival Time and Burst Time for P1 (AT BT): 0 8

Enter Arrival Time and Burst Time for P2 (AT BT): 1 4

Enter Arrival Time and Burst Time for P3 (AT BT): 2 9

Enter Arrival Time and Burst Time for P4 (AT BT): 3 5

===== GANTT CHART =====

```
|P1|=   P2   |=|   P4   |=|   P1   |=|   P3   |=|
0  1       5       10      17      26
```

=====

ProcessID	AT	BT	CT	TAT	(CT-AT)	WT	(TAT-BT)
P1	0	8	17	17	17-0=17	9	17-8=9
P2	1	4	5	4	5-1=4	0	4-4=0
P3	2	9	26	24	26-2=24	15	24-9=15
P4	3	5	10	7	10-3=7	2	7-5=2

Average TAT = 13.00

Average WT = 6.50

Useful CPU time (sum BT) = 26

Total elapsed (from ATmin= 0 to end= 26) = 26

Efficiency (Utilization) = 100.00%

NOTE: The remaining CPU scheduling algorithms such as Priority Scheduling, Round Robin (RR) and ..will be covered in the next lab session.